


Chapter 8: Deadlocks

Traffic deadlock



Chapter Objectives

1. System Model
2. Deadlock in Multithreaded Applications
3. Deadlock **Characterization**
4. Illustrate how deadlock can **occur when mutex locks** are used
5. Define the **four necessary conditions** that characterize deadlock
6. Identify a deadlock situation in a **resource allocation graph**
7. Learn the methods for **handling Deadlocks**

Methods for handling Deadlocks:

1. Evaluate the four different approaches for **preventing deadlocks**
2. Apply the **banker's algorithm** for **deadlock avoidance**
3. Apply the deadlock **detection algorithm**
4. Evaluate approaches for **recovering from deadlock**

Deadlocks happen

Deadlocks happen in **multiprogramming** environment when **several threads** compete for a **finite number of resources**.

A set of threads is in a deadlocked state *when every thread in the set is waiting for an event that can be caused only by another thread in the set*. The events with which we are mainly concerned here are resource acquisition and release.

A thread requests resources; if the resources are **not available** at that time, the process enters a **waiting state**.

- Sometimes, a waiting thread is never again able to change state, because the **resources** it has requested are held by **other waiting threads**.
- This situation is called a **deadlock**.

System Resource Classification

System resources may be partitioned into several types (or **classes**)

Classification:

1. Each class consisting of some number of **identical instances**
2. If a thread requests an instance of a resource type, the allocation of **any instance of the type** should satisfy the request.
3. For example, a system may have two printers.
 1. These two printers may be defined to be in the **same resource class** if no one cares which printer prints which output.
4. **If not**, then the instances are not identical, and **add another class** for it.

Resource Classification

Each class consists of some number of identical instances.

- CPU cycles –
 - ▶ one or more instances: CPU Instance 1, CPU Instance 2, etc.
- Memory
 - ▶ one or more instances: Memory instance 1, Memory instance 2, etc.
- Semaphores
 - ▶ zero or more instances: Semaphore instance 1, instance 2, etc.
- Mutex locks
 - ▶ zero or more instances: Mutex locks instance 1, instance 2, etc.
- Files
 - ▶ zero or more instances: File instance 1, File instance 2 etc.
- I/O devices
 - ▶ zero or more instances: Printer Instance 1, Printer instance 2, etc.

System Model

- ❖ A thread may request as many resources as it requires to carry out its designated task.
- ❖ A thread must request a resource before using it and must release the resource after using it:

1. Request. The thread requests the resource.

- If the request cannot be granted immediately (for example, if the resource is being used by another thread), then the requesting thread must wait until it can acquire the resource.

2. Use.

- The thread can operate on the resource (for example, if the resource is a printer, the thread can print on the printer).

3. Release.

- The thread releases the resource.

System Calls for resource request and release

- device
 - request() and release()
- file
 - open() and close()
- memory
 - allocate() and free()
- semaphores
 - wait() and signal()
- mutex lock
 - acquire() and release()

Deadlock in Multithreaded Applications

- In Figure 8.1 Deadlock Example, `thread_one` attempts to acquire the mutex locks in the order (1) `first_mutex`, (2) `second_mutex`.
- At the same time, `thread_two` attempts to acquire the mutex locks in the order (1) `second_mutex`, (2) `first_mutex`.
- Deadlock is possible if `thread_one` acquires `first_mutex` while `thread_two` acquires `second_mutex`.
- Deadlock will not occur if `thread_one` can acquire and release the mutex locks for `first_mutex` and `second_mutex` before `thread_two` attempts to acquire the locks.
- And, of course, the order in which the threads run depends on how they are scheduled by the CPU scheduler.

Livelock

Livelock is another form of liveness failure.

- Livelock is similar to deadlock. It is less common than deadlock
- Both prevent two or more threads from proceeding, but the threads are unable to proceed for different reasons.
- **Deadlock** occurs when every thread in a set is blocked waiting for an event that can be caused only by another thread in the set.
- **Llivelock** occurs when a thread continuously attempts an action that fails. They aren't blocked, but they aren't making any progress.
- Livelock typically occurs when threads retry failing operations at the same time.
- It thus can generally be avoided by having each thread retry the failing operation at random times.
- This is precisely the approach taken by Ethernet networks when a network collision occurs.

Livelock example

- Livelock can be illustrated with the Pthreads `pthread_mutex_trylock()` function, which attempts to acquire a mutex lock without blocking.
- The code example in Figure 8.2 rewrites the example from Figure 8.1 so that it now uses `pthread_mutex_trylock()`.
- This situation can lead to livelock if thread one acquires first mutex, followed by thread two acquiring second mutex.
- Each thread then invokes `pthread_mutex_trylock()`, which fails, releases their respective locks, and repeats the same actions indefinitely.

Deadlock

Livelock

```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}

/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

Figure 8.1 Deadlock example.

```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    int done = 0;

    while (!done) {
        pthread_mutex_lock(&first_mutex);
        if (pthread_mutex_trylock(&second_mutex)) {
            /**
             * Do some work
             */
            pthread_mutex_unlock(&second_mutex);
            pthread_mutex_unlock(&first_mutex);
            done = 1;
        }
        else
            pthread_mutex_unlock(&first_mutex);
    }

    pthread_exit(0);
}

/* thread_two runs in this function */
void *do_work_two(void *param)
{
    int done = 0;

    while (!done) {
        pthread_mutex_lock(&second_mutex);
        if (pthread_mutex_trylock(&first_mutex)) {
            /**
             * Do some work
             */
            pthread_mutex_unlock(&first_mutex);
            pthread_mutex_unlock(&second_mutex);
            done = 1;
        }
        else
            pthread_mutex_unlock(&second_mutex);
    }

    pthread_exit(0);
}
```

Figure 8.2 Livelock example.

Necessary Conditions

Deadlock can arise if **four conditions hold simultaneously**.

- **Mutual exclusion:** At least one resource must be held in a nonsharable mode; that is, only **one** thread **at a time** can use **a resource**. If another thread requests that resource, the requesting thread must be delayed until the resource has been released.
- 1. **Hold and wait:** A thread must be **holding at least one** resource and **waiting to acquire additional** resources currently **held by other** threads
- 2. **No preemption:** a resource can be **released only voluntarily** by the thread holding it, **after** that thread has **completed its task**
- 3. **Circular wait:** there exists a set $\{T_0, T_1, \dots, T_n\}$ of waiting threads such that T_0 is **waiting** for a resource that is **held by** T_1 , T_1 is waiting for a resource that is held by T_2 , ..., T_{n-1} is waiting for a resource that is held by T_n , and T_n is waiting for a resource that is held by T_0 .
- ✓ **all four conditions must hold for a deadlock to occur.**

Resource-Allocation Graph


Deadlocks can be described more precisely in terms of a directed graph called a **system resource-allocation graph**.

This graph consists of a set of **vertices V** and a set of **edges E** .

- V is partitioned into two types:
 - $T = \{T_1, T_2, \dots, T_n\}$, the set consisting of all the threads in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- **request edge** – directed edge $T_i \rightarrow R_j$
- **assignment edge** – directed edge $R_j \rightarrow T_i$

Resource-Allocation Graph (Cont.)


- Thread




- Resource Type with 4 instances



- T_i requests instance of R_j



- T_i is holding an instance of R_j



When a thread T_i requests an instance of resource type R_j a **request edge** is inserted in the graph

When request is granted, the request edge is transformed to an **assignment edge**.

When the process no longer needs the resource, it releases the resource and the assignment **edge is deleted**.

Resource Allocation Graph Example

$T = \{T_1, T_2, T_3\}$

$R = \{R_1, R_2, R_3, R_4\}$


$E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

■ Resource instances:

- One instance of R1
- Two instances of R2
- One instance of R3
- Three instance of R4

■ Thread states:

- T1 holds one instance of R2 and is waiting for an instance of R1
- T2 holds one instance of R1, one instance of R2, and is waiting for an instance of R3
- T3 holds one instance of R3



Cycle in a resource allocation graph

1. NO DEADLOCK:

- If the graph contains **no cycles**, then **no** thread in the system **deadlocked**.


2. DEADLOCK:

- If the graph **contains a cycle**, and if the cycle involves only a set of **resource types**, **each** of which has **only a single instance**, then **each thread involved in the cycle is deadlocked**.

3. NOT NECESSARILY:

- If **each resource type** has **several instances**, then a cycle **does not necessarily imply** that a deadlock has occurred. In this case, a cycle in the graph is a necessary but not a sufficient condition for the existence of deadlock.

Resource Allocation Graph With A Deadlock



2 cycles :

$R_2 \rightarrow T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2$

$R_2 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2$


Threads T_1 , T_2 , and T_3 are **deadlocked**.

T_2 is waiting for the resource R_3 , which is held by T_3 .

T_3 is waiting for either thread T_1 or thread T_2 to release resource R_2 .

In addition, T_1 is waiting for T_2 to release resource R_1 .

Graph With A Cycle But No Deadlock



we also have a cycle:
 $T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$

There is no deadlock.

Thread T_4 may release its instance of resource type R_2 .

That resource can then be allocated to T_3 , breaking the cycle.

Basic Facts

- If graph contains **no cycles** \Rightarrow no deadlock
- If graph contains a **cycle** \Rightarrow
 - if only **one instance per resource type**, then deadlock
 - if **several instances per resource type**, possibility of deadlock

Methods for Handling Deadlocks

- ❑ Ensure that the system will **never** enter a deadlock state:
 - ❑ Deadlock **prevention**
 - ❑ Deadlock **avoidance**
- ❑ **Allow** the system to enter a deadlock state and then **recover**
 - solution **used by most operating systems**, including **Linux and Windows**.
 - It is then up to the application developer to write programs that handle deadlocks.
- ❑ **Ignore** the problem and pretend that deadlocks never occur in the system.

1. Deadlock Prevention

Invalidate **one of the four necessary conditions** for deadlock:

- **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); **must hold for non-sharable resources**
 - **cannot prevent deadlocks by denying the mutual-exclusion condition**, because **some resources** are intrinsically **nonsharable**.
- **Hold and Wait** – must guarantee that whenever a thread **requests a resource, it does not hold** any other resources
 - Require a thread to **request and be allocated** all its resources **before it begins execution**, or allow thread to request resources only when the thread has none allocated to it.
 - Low resource utilization; starvation possible

Deadlock Prevention (Cont.)

■ No Preemption –

- If a thread that is holding some resources requests another resource that **cannot be** immediately **allocated** to it, then **all resources currently being held are released**
- Preempted resources are added to the list of resources for which the thread is waiting
- Thread will be **restarted only when it can regain its old resources, as well as the new ones** that it is requesting

■ Circular Wait – **impose a total ordering** of all resource types, and **require** that each thread requests resources in an **increasing order** of enumeration

Circular Wait

- Invalidating the circular wait condition is most common.
- Simply **assign each resource (i.e. mutex locks) a unique number.**
- Resources must be acquired in order.
- If:

```
first_mutex = 1
second_mutex = 5
```

code for `thread_two` could not be written as follows:

Must follow the **increasing order**:

Request `first_mutex`,
then `second_mutex`

```
/* thread.one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}

/* thread.two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

2. Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that **each thread declare the *maximum number of resources of each type that it may need***
- The deadlock-avoidance algorithm **dynamically examines the resource-allocation state** to ensure that there **can never be a circular-wait condition**
- Resource-allocation **state** is defined by the **number of available and allocated resources, and the maximum demands of the threads**

Safe State

A **state is safe** if the system can allocate resources to each thread (up to its maximum) in some order and still avoid a deadlock.

A system is in a safe state only if there exists a **safe sequence**.


A sequence of threads $\langle T_1, T_2, \dots, T_n \rangle$ is a safe sequence for the current allocation state if, for each T_i , the resource requests that T_i can still make, can be satisfied by the currently available resources plus the resources held by all T_j , with $j < i$.

That is:

- If the resources that T_i needs are not immediately available, then T_i can wait until all T_j have finished.
- When they have finished, T_i can obtain all of its needed resources, complete its designated task, return its allocated resources, and terminate.
- When T_i terminates, T_{i+1} can obtain its needed resources, and so on.
- **If no such sequence exists**, then the system state is said to be **unsafe**.

Safe, Unsafe, Deadlock State

- A safe state is not a deadlocked state.
- Conversely, a deadlocked state is an unsafe state.
- Not all unsafe states are deadlocks
- An unsafe state **may** lead to a deadlock.
- As long as the state is safe, the operating system can avoid unsafe (and deadlocked) states.
- In an unsafe state, the operating system cannot prevent threads from requesting resources in such away that a deadlock occurs. The behavior of the threads controls unsafe states.



Safe State illustration

Consider a system with **twelve resources** and **three threads** T_0 , T_1 , and T_2 .

Thread T_0 requires ten resources, thread T_1 may need four, and thread T_2 may need nine resources.

Suppose that, at **time t_0** , thread T_0 is holding **five** resources, thread T_1 is holding **two** resources, and thread T_2 is holding **two** resources. (Thus, there are **three free** resources.).

	<u>Maximum Needs</u>	<u>Allocation</u>	<u>Need</u>
T_0	10	5	5
T_1	4	2	2
T_2	9	2	7

At time t_0 , the system is in a safe state. The **sequence $\langle T_1, T_0, T_2 \rangle$ satisfies the safety condition.**

- Thread T_1 can immediately be allocated all its resources and then return them (the system will then have five available resources);
- Then thread T_0 can get all its resources and return them (the system will then have ten available resources);
- Finally thread T_2 can get all its resources and return them (the system will then have all twelve resources available).

Basic Facts

- A system can go from a safe state to an unsafe state.
- Suppose that, at time t_1 , thread T_2 requests and is allocated one more resource. Only thread T_1 can be allocated all its resources. When it returns them, the system will have only four available resources. The system is no longer in a safe state.
- If a system is in **safe state** \Rightarrow **no deadlocks**
- If a system is in **unsafe state** \Rightarrow **possibility of deadlock**
- **Avoidance** \Rightarrow **ensure** that a system will **never** enter an **unsafe state**.

Avoidance Algorithms


- Single instance of a resource type
 - Use a resource-allocation graph
- Multiple instances of a resource type
 - Use the Banker's Algorithm

Resource-Allocation Graph Scheme

In addition to the request and assignment edges, we introduce a new type of edge, called a **claim edge**, represented by a **dashed line**


- **Claim edge** $T_i \rightarrow R_j$ indicated that thread T_i may request resource R_j at **some time in the future**.
- Claim edge converts to request edge when a thread requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the thread
- When a resource is released by a thread, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

Resource-Allocation Graph for Deadlock Avoidance



Now suppose that T_i requests resource R_j . The request can be **granted only** if converting the request edge $T_i \rightarrow R_j$ to an assignment edge $R_j \rightarrow T_i$ **does not result** in the formation of a **cycle** in the graph

Unsafe State In Resource-Allocation Graph



Although R_2 is currently free, we cannot allocate it to T_2 , since this action will **create a cycle** in the graph. A cycle indicates that the system is in an **unsafe state**.

Resource-Allocation Graph Algorithm

- Suppose that thread T_i requests a resource R_j
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

Banker's Algorithm

- Multiple instances of resources
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If $Available[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Example of Banker's Algorithm

- 5 processes P_0 through P_4

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$		$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2	P_0	7 4 3
P_1	2 0 0	3 2 2		P_1	1 2 2
P_2	3 0 2	9 0 2		P_2	6 0 0
P_3	2 1 1	2 2 2		P_3	0 1 1
P_4	0 0 2	4 3 3		P_4	4 3 1

Need is Max – Allocation

$7-0 = 7$ $5-1=4$ $3-0=3$

$3-2 = 1$ $2-0=2$ $2-0=2$

$9-3 = 6$ $0-0=0$ $2-2=0$

$2-2 = 0$ $2-2=1$ $2-1=1$

$4-0 = 4$ $3-0=3$ $3-2=1$

Example (Cont.)

- The content of the matrix **Need** is defined to be **Max – Allocation**

	<u>Need</u>
	A B C
P_0	7 4 3
P_1	1 2 2
P_2	6 0 0
P_3	0 1 1
P_4	4 3 1

- The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria

Example: P_1 Request (1,0,2)

- Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow$ true

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>
	A	B	C	A	B	C	A B C
P_0	0	1	0	7	4	3	2 3 0
P_1			3 0 2			0 2 0	
P_2	3	0	2	6	0	0	
P_3	2	1	1	0	1	1	
P_4	0	0	2	4	3	1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement
- Can request for (3,3,0) by P_4 be granted?
 - No. Resources are not available
- Can request for (0,2,0) by P_0 be granted?
 - No. Resources available, but will result in unsafe state

3. Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Deadlock detection algorithm – Single Instance


- Maintain **wait-for** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically invoke the algorithm that searches for a **cycle in the graph**. If there is a cycle, there exists a deadlock

If all resources have only a single instance, define a deadlock detection algorithm that uses a variant of the resource-allocation graph, called a **wait-for** graph.

We obtain this graph from the resource-allocation graph by removing the resource nodes and collapsing the appropriate edges.


- More precisely, an edge from P_i to P_j in a wait-for graph implies that process P_i is waiting for process P_j to release a resource that P_i needs.
- (An edge $P_i \rightarrow P_j$ exists in a wait-for graph if and only if the corresponding resource allocation graph contains two edges $P_i \rightarrow Rq$ and $Rq \rightarrow P_j$ for some resource Rq)

Resource-Allocation Graph and Wait-for Graph



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph

Deadlock detection algorithm - Several Instances of a Resource Type

- **Available:** A vector of length m indicates the number of available resources of each type
- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process
- **Request:** An $n \times m$ matrix indicates the current request of each process. If **Request** $[i][j] = k$, then process P_i is requesting k more instances of resource type R_j .

Detection Algorithm

1. Let **Work** and **Finish** be vectors of length **m** and **n**, respectively
Initialize:
 - (a) **Work = Available**
 - (b) For $i = 1, 2, \dots, n$, if **Allocation_i $\neq 0$** , then
Finish[i] = false; otherwise, **Finish[i] = true**
2. Find an index **i** such that both:
 - (a) **Finish[i] == false**
 - (b) **Request_i \leq Work**

If no such **i** exists, go to step 4

Detection Algorithm (Cont.)

3. **$Work = Work + Allocation_i$**
 $Finish[i] = true$
go to step 2
4. If **$Finish[i] == false$** , for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if **$Finish[i] == false$** , then P_i is deadlocked

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state

Example of Detection Algorithm

- Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances)

- Snapshot at time T_0 :

	<u>Allocation</u>			<u>Request</u> <u>Available</u>		
	A	B	C	A	B	C
P_0	0	1	0	0	0	0
P_1	2	0	0	2	0	2
P_2			3	0	3	
P_3	2	1	1	1	0	0
P_4	0	0	2	0	0	2

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in ***Finish[i] = true*** for all i

Example (Cont.)

- P_2 requests an additional instance of type **C**

	<u>Request</u>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_0	0	0	0
P_1	2	0	2
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- State of system?
 - Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes; requests
 - Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Detection-Algorithm Usage

- When, and how often, to invoke depends on:
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - ▶ one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
 1. Priority of the process
 2. How long process has computed, and how much longer to completion
 3. Resources the process has used
 4. Resources process needs to complete
 5. How many processes will need to be terminated
 6. Is process interactive or batch?

4. Recovery from Deadlock: Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor

End of Chapter 8